

- Omar Shaikh, Valentino Chai, Michele J Gelfand, Diyi Yang, and Michael S Bernstein. Rehearsal: Simulating conflict to teach conflict resolution. *ArXiv preprint*, abs/2309.12309, 2023. URL <https://arxiv.org/abs/2309.12309>.
- Omar Shaikh, Valentino Emil Chai, Michele Gelfand, Diyi Yang, and Michael S Bernstein. Rehearsal: Simulating conflict to teach conflict resolution. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pp. 1–20, 2024.
- Murray Shanahan, Kyle McDonell, and Laria Reynolds. Role play with large language models. *Nature*, 2023. doi: 10.1038/s41586-023-06647-8. URL <https://doi.org/10.1038/s41586-023-06647-8>.
- Tianhao Shen, Renren Jin, Yufei Huang, Chuang Liu, Weilong Dong, Zishan Guo, Xinwei Wu, Yan Liu, and Deyi Xiong. Large language model alignment: A survey. *ArXiv preprint*, abs/2309.15025, 2023. URL <https://arxiv.org/abs/2309.15025>.
- Zijing Shi, Meng Fang, Shunfeng Zheng, Shilong Deng, Ling Chen, and Yali Du. Cooperation on the fly: Exploring language agents for ad hoc teamwork in the avalon game. *ArXiv preprint*, abs/2312.17515, 2023. URL <https://arxiv.org/abs/2312.17515>.
- Petter Törnberg, Diliara Valeeva, Justus Uitermark, and Christopher Bail. Simulating social media using large language models to evaluate alternative news feed algorithms. *ArXiv preprint*, abs/2310.05984, 2023. URL <https://arxiv.org/abs/2310.05984>.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *ArXiv preprint*, abs/2307.09288, 2023. URL <https://arxiv.org/abs/2307.09288>.
- Maximilian Puelma Touzel, Sneheel Sarangi, Austin Welch, Gayatri Krishnakumar, Dan Zhao, Zachary Yang, Hao Yu, Ethan Kosak-Hine, Tom Gibbs, Andreea Musulan, et al. A simulation system towards solving societal-scale manipulation. *arXiv preprint arXiv:2410.13915*, 2024.
- Milena Tsvetkova, Taha Yasseri, Niccolo Pescetelli, and Tobias Werner. A new sociology of humans and machines. *Nature Human Behaviour*, 8(10):1864–1876, 2024.
- Eric M Uslaner. Producing and consuming trust. *Political science quarterly*, 115(4):569–590, 2000.
- Lei Wang, Jingsen Zhang, Xu Chen, Yankai Lin, Ruihua Song, Wayne Xin Zhao, and Ji-Rong Wen. Recagent: A novel simulation paradigm for recommender systems. *ArXiv preprint*, abs/2306.02552, 2023a. URL <https://arxiv.org/abs/2306.02552>.
- Shenzhi Wang, Chang Liu, Zilong Zheng, Siyuan Qi, Shuo Chen, Qisen Yang, Andrew Zhao, Chaofei Wang, Shiji Song, and Gao Huang. Avalon’s game of thoughts: Battle against deception through recursive contemplation. *ArXiv preprint*, abs/2310.01320, 2023b. URL <https://arxiv.org/abs/2310.01320>.
- Yufei Wang, Wanjuan Zhong, Liangyou Li, Fei Mi, Xingshan Zeng, Wenyong Huang, Lifeng Shang, Xin Jiang, and Qun Liu. Aligning large language models with human: A survey. *ArXiv preprint*, abs/2307.12966, 2023c. URL <https://arxiv.org/abs/2307.12966>.
- Oliver E Williamson. Calculativeness, trust, and economic organization. *The journal of law and economics*, 36(1, Part 2):453–486, 1993.
- Ruoxi Xu, Yingfei Sun, Mengjie Ren, Shiguang Guo, Ruotong Pan, Hongyu Lin, Le Sun, and Xianpei Han. Ai for social science and social science of ai: A survey. *arXiv preprint arXiv: 2401.11839*, 2024.
- Yuzhuang Xu, Shuo Wang, Peng Li, Fuwen Luo, Xiaolong Wang, Weidong Liu, and Yang Liu. Exploring large language models for communication games: An empirical study on werewolf. *ArXiv preprint*, abs/2309.04658, 2023. URL <https://arxiv.org/abs/2309.04658>.
- Diyi Yang, Caleb Ziems, William Held, Omar Shaikh, Michael S Bernstein, and John Mitchell. Social skill training with large language models. *ArXiv preprint*, abs/2404.04204, 2024. URL <https://arxiv.org/abs/2404.04204>.

- Murong Yue, Wijidane Mifdal, Yixuan Zhang, Jennifer Suh, and Ziyu Yao. Mathvc: An llm-simulated multi-character virtual classroom for mathematics education. *ArXiv preprint*, abs/2404.06711, 2024. URL <https://arxiv.org/abs/2404.06711>.
- An Zhang, Leheng Sheng, Yuxin Chen, Hao Li, Yang Deng, Xiang Wang, and Tat-Seng Chua. On generative agents in recommendation. *ArXiv preprint*, abs/2310.10108, 2023a. URL <https://arxiv.org/abs/2310.10108>.
- Jintian Zhang, Xin Xu, and Shumin Deng. Exploring collaboration mechanisms for llm agents: A social psychology view. *ArXiv preprint*, abs/2310.02124, 2023b. URL <https://arxiv.org/abs/2310.02124>.
- Xinnong Zhang, Jiayu Lin, Libo Sun, Weihong Qi, Yihang Yang, Yue Chen, Hanjia Lyu, Xinyi Mou, Siming Chen, Jiebo Luo, Xuanjing Huang, Shiping Tang, and Zhongyu Wei. Electionsim: Massive population election simulation powered by large language model driven agents. *arXiv preprint arXiv: 2410.20746*, 2024. URL <https://arxiv.org/abs/2410.20746>.
- Xuhui Zhou, Hao Zhu, Leena Mathur, Ruohong Zhang, Haofei Yu, Zhengyang Qi, Louis-Philippe Morency, Yonatan Bisk, Daniel Fried, Graham Neubig, et al. Sotopia: Interactive evaluation for social intelligence in language agents. *ArXiv preprint*, abs/2310.11667, 2023. URL <https://arxiv.org/abs/2310.11667>.
- Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. Can large language models transform computational social science? *ArXiv preprint*, abs/2305.03514, 2023. URL <https://arxiv.org/abs/2305.03514>.

Content of Appendix

A	Related Work	18
B	Impact Statement	18
C	Limitations and Future Works	19
D	Additional Illustration for Experiments on Risk Perception	19
E	Statistical Testing	20
F	More Experiments on Probing Intrinsic Properties of Agent Trust	21
G	The Complete Results for the Repeated Trust Game	22
G.1	Human	22
G.2	GPT-4	23
G.3	GPT-3.5	24
H	Prompt Setting	25
H.1	Persona Prompt	25
H.2	Game Setting Prompt	26
H.3	Prompts for Probing Intrinsic Properties	28
I	Belief-Desire-Intention (BDI) Analysis	30
I.1	GPT-4 in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent)	30
I.2	GPT-3.5-turbo-0613 in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent)	31
I.3	text-davinci-003 in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent) . .	32
I.4	GPT-3.5-turbo-instruct in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent)	33
I.5	Llama2-13b in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent)	34
I.6	Llama2-70b in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent)	35
I.7	Vicuna-v1.3-7b in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent) . . .	36
I.8	Vicuna-v1.3-13b in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent) . .	37
I.9	Vicuna-v1.3-33b in the Trust Game (<i>Low</i> Amount Sent vs. <i>High</i> Amount Sent) . .	38
I.10	the Dictator Game vs. the Trust Game	39
I.11	the MAP Trust Game	40
I.12	the Lottery Game	41
I.13	the Repeated Trust Game	42
I.14	the Trust Game + Gender	46
I.15	the Trust Game + <i>Agents</i> vs. <i>Human</i>	47
I.16	the Trust Game + Trust Manipulation	48
I.17	the Trust Game + No CoT vs CoT	49