

- 
- Cheng, J., Marone, M., Weller, O., Lawrie, D., Khashabi, D., & Van Durme, B. (2024). Dated data: Tracing knowledge cutoffs in large language models. *arXiv:2403.12958*. <https://doi.org/10.48550/arXiv.2403.12958>
- Chuang, Y.-S., Studdiford, Z., Nirunwiroj, K., Goyal, A., Frigo, V. V., Yang, S., . . . Rogers, T. T. (2024). Beyond demographics: Aligning role-playing LLM-based agents using human belief networks. *arXiv:2406.17232*. <https://doi.org/10.48550/arXiv.2406.17232>
- Church, A. (1936). An unsolvable problem of elementary number theory. *American Journal of Mathematics*, 58(2), 345-363. <https://doi.org/10.2307/2371045>
- Colombatto, C., & Fleming, S. M. (2024). Folk psychological attributions of consciousness to large language models. *Neuroscience of Consciousness*, 2024(1), niae013. <https://doi.org/10.1093/nc/naie013>
- Cook, T. D., & Campbell, D. T. (1979). *Quasi-experimentation: Design & analysis issues for field settings*. Houghton Mifflin.
- Crockett, M., & Messeri, L. (2023). Should large language models replace human participants? *PsyArXiv*. <https://doi.org/10.31234/osf.io/4zdx9>
- Cronbach, L. J., & Meehl, P. E. (1955). Construct validity in psychological tests. *Psychological Bulletin*, 52(4), 281-302. <https://doi.org/10.1037/h0040957>
- Dasgupta, I., Lampinen, A. K., Chan, S. C., Sheahan, H. R., Creswell, A., Kumaran, D., . . . Hill, F. (2022). Language models show human-like content effects on reasoning tasks. *arXiv:2207.07051*. <https://doi.org/10.48550/arXiv.2207.07051>
- Demszky, D., Yang, D., Yeager, D. S., Bryan, C. J., Clapper, M., Chandhok, S., . . . Pennebaker, J. W. (2023). Using large language models in psychology. *Nature Reviews Psychology*, 2(11), 688-701. <https://doi.org/10.1038/s44159-023-00241-5>
- Dettki, H. M., Lake, B. M., Wu, C. M., & Rehder, B. (2025). Do large language models reason causally like us? Even better? *arXiv:2502.10215*. <https://doi.org/10.48550/arXiv.2502.10215>
- Dillion, D., Tandon, N., Gu, Y., & Gray, K. (2023). Can AI language models replace human participants? *Trends in Cognitive Sciences*, 27(7), 597-600. <https://doi.org/10.1016/j.tics.2023.04.008>
- Doerig, A., Sommers, R. P., Seeliger, K., Richards, B., Ismael, J., Lindsay, G. W., . . . Kietzmann, T. C. (2023). The neuroconnectionist research programme. *Nature Reviews Neuroscience*, 24(7), 431-450. <https://doi.org/10.1038/s41583-023-00705-w>
- Ferreira, F., Bailey, K. G. D., & Ferraro, V. (2002). Good-enough representations in language comprehension. *Current Directions in Psychological Science*, 11(1), 11-15. <https://doi.org/10.1111/1467-8721.00158>
- Firestone, C. (2020). Performance vs. competence in human-machine comparisons. *Proceedings of the National Academy of Sciences of the United States of America*, 117(43), 26562-26571. <https://doi.org/10.1073/pnas.1905334117>
- Fokas, A. S. (2023). Can artificial intelligence reach human thought? *PNAS Nexus*, 2(12), pgad409. <https://doi.org/10.1093/pnasnexus/pgad409>

- Gao, Y., Lee, D., Burtch, G., & Fazelpour, S. (2024). Take caution in using LLMs as human surrogates: Scylla ex machina. *arXiv:2410.19599*.  
<https://doi.org/10.48550/arXiv.2410.19599>
- Grossmann, I., Feinberg, M., Parker, D. C., Christakis, N. A., Tetlock, P. E., & Cunningham, W. A. (2023). AI and the transformation of social science research. *Science*, 380(6650), 1108-1109. <https://doi.org/10.1126/science.ad1778>
- Guest, O., & Martin, A. E. (2023). On logical inference over brains, behaviour, and artificial neural networks. *Computational Brain & Behavior*, 6(2), 213-227.  
<https://doi.org/10.1007/s42113-022-00166-x>
- Gui, G., & Toubia, O. (2023). The challenge of using LLMs to simulate human behavior: A causal inference perspective. *arXiv:2312.15524*.  
<https://doi.org/10.48550/arXiv.2312.15524>
- Gurnee, W., & Tegmark, M. (2023). Language models represent space and time. *arXiv:2310.02207*. <https://doi.org/10.48550/arXiv.2310.02207>
- GX-Chen, A., Lin, D., Samiei, M., Precup, D., Richards, B. A., Fergus, R., & Marino, K. (2025). Language agents mirror human causal reasoning biases: How can we help them think like scientists? *arXiv:2505.09614*.  
<https://doi.org/10.48550/arXiv.2505.09614>
- Hagendorff, T., Fabi, S., & Kosinski, M. (2023). Human-like intuitive behavior and reasoning biases emerged in large language models but disappeared in ChatGPT. *Nature Computational Science*, 3(10), 833-838. <https://doi.org/10.1038/s43588-023-00527-x>
- Harding, J., D'Alessandro, W., Laskowski, N. G., & Long, R. (2024). AI language models cannot replace human research participants. *AI & Society*, 39, 2603–2605.  
<https://doi.org/10.1007/s00146-023-01725-x>
- Harding, J., & Sharadin, N. (in press). What is it for a machine learning model to have a capability? *The British Journal for the Philosophy of Science*.  
<https://doi.org/10.1086/732153>
- Hu, J., Mahowald, K., Lupyan, G., Ivanova, A., & Levy, R. (2024). Language models align with human judgments on key grammatical constructions. *Proceedings of the National Academy of Sciences of the United States of America*, 121(36), e2400917121. <https://doi.org/10.1073/pnas.2400917121>
- Hu, T., Kyrychenko, Y., Rathje, S., Collier, N., van der Linden, S., & Roozenbeek, J. (2025). Generative language models exhibit social identity biases. *Nature Computational Science*, 5(1), 65-75. <https://doi.org/10.1038/s43588-024-00741-1>
- Huang, P.-H., Lin, Z., Imbot, S., Fu, W., & Tu, E. (2025). Analysis of LLM bias (Chinese propaganda and anti-US sentiment) in DeepSeek-R1 vs. ChatGPT o3-mini-high. *arXiv:2506.01814*. <https://doi.org/10.48550/arXiv.2506.01814>
- Ibrahim, L., & Cheng, M. (2025). Thinking beyond the anthropomorphic paradigm benefits LLM research. *arXiv:2502.09192*. <https://doi.org/10.48550/arXiv.2502.09192>
- Ivanova, A. A. (2025). How to evaluate the cognitive abilities of LLMs. *Nature Human Behaviour*, 9(2), 230-233. <https://doi.org/10.1038/s41562-024-02096-z>

- Jackson, F. (1982). Epiphenomenal qualia. *The Philosophical Quarterly*, 32(127), 127-136.  
<https://doi.org/10.2307/2960077>
- Jackson, F. (1986). What Mary didn't know. *The Journal of Philosophy*, 83(5), 291-295.  
<https://doi.org/10.2307/2026143>
- Jin, Z., Kleiman-Weiner, M., Piatti, G., Levine, S., Liu, J., Gonzalez, F., . . . Mihalcea, R. (2024). Language model alignment in multilingual trolley problems. *arXiv:2407.02273*. <https://doi.org/10.48550/arXiv.2407.02273>
- Jones, C. R., & Bergen, B. (2024). Does word knowledge account for the effect of world knowledge on pronoun interpretation? *Language and Cognition*, 1-32.  
<https://doi.org/10.1017/langcog.2024.2>
- Jones, C. R., Chang, T. A., Coulson, S., Michaelov, J. A., Trott, S., & Bergen, B. (2022). Distributional semantics still can't account for affordances. Proceedings of the Annual Meeting of the Cognitive Science Society, Retrieved from  
<https://escholarship.org/uc/item/44z7r3j3>.
- Kamoi, R., Das, S. S. S., Lou, R., Ahn, J. J., Zhao, Y., Lu, X., . . . Vummanthala, S. R. (2024). Evaluating LLMs at detecting errors in LLM responses. *arXiv:2404.03602*.  
<https://doi.org/10.48550/arXiv.2404.03602>
- Kerschbaumer, S., Voracek, M., Aczél, B., Anderson, S. F., Booth, B. M., Buchanan, E. M., . . . Tran, U. S. (2025). VALID: A checklist-based approach for improving validity in psychological research. *Advances in Methods and Practices in Psychological Science*, 8(1), 25152459241306432.  
<https://doi.org/10.1177/25152459241306432>
- Kim, J. S., Elli, G. V., & Bedny, M. (2019). Knowledge of animal appearance among sighted and blind adults. *Proceedings of the National Academy of Sciences of the United States of America*, 116(23), 11213-11222. <https://doi.org/10.1073/pnas.1900952116>
- Kozlowski, A. C., & Evans, J. A. (2024). Simulating subjects: The promise and peril of AI stand-ins for social agents and interactions. *SocArXiv*.  
<https://doi.org/10.31235/osf.io/vp3j2>
- Lahoti, P., Blumm, N., Ma, X., Kotikalapudi, R., Potluri, S., Tan, Q., . . . Beutel, A. (2023). Improving diversity of demographic representation in large language models via collective-critiques and self-voting. *arXiv:2310.16523*.  
<https://doi.org/10.48550/arXiv.2310.16523>
- Lake, B. M., & Murphy, G. L. (2023). Word meaning in minds and machines. *Psychological Review*, 130(2), 401-431. <https://doi.org/10.1037/rev0000297>
- Lee, M. H., Montgomery, J. M., & Lai, C. K. (2024). Large language models portray socially subordinate groups as more homogeneous, consistent with a bias observed in humans. Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency, Rio de Janeiro, Brazil.
- Lee, S., Li, M., Lai, B., Jia, W., Ryan, F., Cao, X., . . . Yang, D. (2024). Towards social AI: A survey on understanding social interactions. *arXiv:2409.15316*.  
<https://doi.org/10.48550/arXiv.2409.15316>